

Free API Keys — Sources & What They're For

Create your own free account on each site and use your own key, within the free quota. Free tier ≠ stolen key.

🧠 LLM / "Brain" APIs

What it does: takes text in, returns generated text. **Where to use:** chatbots & voice-assistant brains (like Jarvis), tool-calling agents, summarizing/rewriting, classification & extraction, content & marketing copy, coding help, Q&A over your data.

PROVIDER	GET THE KEY	BEST USED FOR
Google Gemini	aistudio.google.com/apikey	General brain + vision (images). Best free tier. Already wired in Jarvis (AIza...)
Groq	console.groq.com	When you need speed — near-instant replies for live voice/chat (Llama, Qwen, gpt-oss)
Cerebras	cloud.cerebras.ai	Also very fast inference; good real-time fallback
OpenRouter	openrouter.ai/keys	Try many models with one key ; some free. Good for experimenting
Mistral	console.mistral.ai	Lightweight chat / coding tasks
Cohere	dashboard.cohere.com	Search/RAG, embeddings, text classification
Together AI	api.together.ai	Wide open-model catalog; signup credits
GitHub Models	github.com/marketplace/models	Free for devs to prototype (uses your GitHub token)
Cloudflare Workers AI	dash.cloudflare.com	Run models at the edge inside web apps
Ollama (local)	ollama.com	Private/offline brain — no key, no internet, unlimited

For the Jarvis fallback: Gemini is the easiest drop-in (already supported). Groq / Cerebras / OpenRouter use the OpenAI format and need ~30 lines added to `brain.ts`.

🗣️ Voice / TTS APIs

What it does: turns text into spoken audio. **Where to use:** the voice that Jarvis speaks with, screen readers / accessibility, audio notifications, IVR phone menus, narration for videos.

PROVIDER	WHERE	BEST USED FOR
Piper	local (installed)	Offline, unlimited voice — your default; works with no internet
macOS say	built-in	Zero-setup fallback so voice never fully drops
Azure Speech	azure.microsoft.com	High-quality neural voices, 500k chars/mo free
Google Cloud TTS	cloud.google.com/text-to-speech	Natural voices, many languages
Amazon Polly	aws.amazon.com/polly	Production TTS (free 12 months)
ElevenLabs	elevenlabs.io	Most realistic "butler" voice; 10k chars/mo free (resets monthly)

🔍 Other Useful APIs

What it does: fetch live real-world data your assistant/app can act on. **Where to use:** answering questions and powering tools.

TYPE	PROVIDER	USED FOR
Weather	Open-Meteo (no key) · OpenWeatherMap	"What's the weather?" — current & forecast data
Web search	DuckDuckGo · Brave · Tavily	Look up live info, feed results to the LLM (RAG)
Flights	Amadeus (test tier)	Search real flights between airports/dates
Maps/places	OpenStreetMap / Nominatim (no key)	Geocoding, addresses, nearby places

Where each fits in a project

Understand & reply to the user	→ LLM / Brain API (Gemini, Groq, Ollama...)
Speak the reply out loud	→ TTS API (Piper, ElevenLabs, Azure...)
Hear the user (speech → text)	→ STT (Whisper — runs locally in Jarvis)
Fetch live facts (weather, news, flights)	→ Data APIs (Open-Meteo, search, Amadeus...)
Answer over your own docs/data	→ Embeddings + search (Cohere, OpenRouter...)

The one rule that keeps it legal

- Sign up with **your own account**, use **your own key**, stay within the **free quota**, and follow each provider's **terms** (usually one free account per person, no reselling, no abuse).
- Never use keys you found posted online or that belong to someone else.

Generated for Arju Jarvis · free tiers & quotas change — check each provider's pricing page for current limits.